

Project name - Census_income_data

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore') # importing python Libraries
```

In [2]:

```
df = pd.read_csv("C:/Users/harshitagups/Desktop/project/census_income_data.csv") # importing dataset
```

In [3]:

```
print('Rows: {} Columns: {}'.format(df.shape[0], df.shape[1])) # defines (rows,columns)
```

Rows: 32561 Columns: 15

In [4]:

```
df.head(11) #first 10 data entry from dataset
```

Out[4]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.g
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	
7	74	State-gov	88638	Doctorate	16	Never-married	Prof-specialty	Other-relative	White	Female	
8	68	Federal-gov	422013	HS-grad	9	Divorced	Prof-specialty	Not-in-family	White	Female	
9	41	Private	70037	Some-college	10	Never-married	Craft-repair	Unmarried	White	Male	
10	45	Private	172274	Doctorate	16	Divorced	Prof-specialty	Unmarried	Black	Female	

In [5]:

```
df.tail(15)
```

#Last 15 data entry from data

Out[5]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	ca
32546	31	Private	199655	Masters	14	Divorced	Other-service	Not-in-family	Other	Female	
32547	39	Local-gov	111499	Assoc-acdm	12	Married-civ-spouse	Adm-clerical	Wife	White	Female	
32548	37	Private	198216	Assoc-acdm	12	Divorced	Tech-support	Not-in-family	White	Female	
32549	43	Private	260761	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
32550	43	State-gov	255835	Some-college	10	Divorced	Adm-clerical	Other-relative	White	Female	
32551	43	Self-emp-not-inc	27242	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	Male	
32552	32	Private	34066	10th	6	Married-civ-spouse	Handlers-cleaners	Husband	Amer-Indian-Eskimo	Male	
32553	43	Private	84661	Assoc-voc	11	Married-civ-spouse	Sales	Husband	White	Male	
32554	32	Private	116138	Masters	14	Never-married	Tech-support	Not-in-family	Asian-Pac-Islander	Male	
32555	53	Private	321865	Masters	14	Married-civ-spouse	Exec-managerial	Husband	White	Male	
32556	22	Private	310152	Some-college	10	Never-married	Protective-serv	Not-in-family	White	Male	
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	
32558	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
32559	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	
32560	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	

In [6]:

```
df.info()                                     #all information regarding dataset like datatypes,
#Observations:
#1. There are in total 32561 samples in the census_income data set
#2. There are both categorical and numerical attributes in the dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   32561 non-null  int64
 1   workclass              32561 non-null  object
 2   fnlwgt                32561 non-null  int64
 3   education              32561 non-null  object
 4   education.num          32561 non-null  int64
 5   marital.status         32561 non-null  object
 6   occupation             32561 non-null  object
 7   relationship           32561 non-null  object
 8   race                   32561 non-null  object
 9   sex                    32561 non-null  object
10   capital.gain           32561 non-null  int64
11   capital.loss           32561 non-null  int64
12   hours.per.week         32561 non-null  int64
13   native.country         32561 non-null  object
14   income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

In [7]:

```
df.nunique()                                     #finding out no. of unique values in part
```

Out[7]:

```
age                73
workclass           9
fnlwgt            21648
education           16
education.num       16
marital.status       7
occupation          15
relationship         6
race                 5
sex                  2
capital.gain        119
capital.loss         92
hours.per.week       94
native.country       42
income              2
dtype: int64
```

In [8]:

```
for i, col in enumerate(df.columns):  
    print(df.columns[i],":", df[str(col)].unique(), '\n')  
#Unique values in c
```

```
age : [90 82 66 54 41 34 38 74 68 45 52 32 51 46 57 22 37 29 61 21 33 49 23 59  
60 63 53 44 43 71 48 73 67 40 50 42 39 55 47 31 58 62 36 72 78 83 26 70  
27 35 81 65 25 28 56 69 20 30 24 64 75 19 77 80 18 17 76 79 88 84 85 86  
87]  
  
workclass : ['?' 'Private' 'State-gov' 'Federal-gov' 'Self-emp-not-inc' 'Self-emp-inc'  
'Local-gov' 'Without-pay' 'Never-worked']  
  
fnlwgt : [ 77053 132870 186061 ... 34066 84661 257302]  
  
education : ['HS-grad' 'Some-college' '7th-8th' '10th' 'Doctorate' 'Prof-school'  
'Bachelors' 'Masters' '11th' 'Assoc-acdm' 'Assoc-voc' '1st-4th' '5th-6th'  
'12th' '9th' 'Preschool']  
  
education.num : [ 9 10 4 6 16 15 13 14 7 12 11 2 3 8 5 1]  
  
marital.status : ['Widowed' 'Divorced' 'Separated' 'Never-married' 'Married-civ-spouse'  
'Married-spouse-absent' 'Married-AF-spouse']  
  
occupation : ['?' 'Exec-managerial' 'Machine-op-inspct' 'Prof-specialty'  
'Other-service' 'Adm-clerical' 'Craft-repair' 'Transport-moving'  
'Handlers-cleaners' 'Sales' 'Farming-fishing' 'Tech-support'  
'Protective-serv' 'Armed-Forces' 'Priv-house-serv']  
  
relationship : ['Not-in-family' 'Unmarried' 'Own-child' 'Other-relative' 'Husband' 'Wife']  
  
race : ['White' 'Black' 'Asian-Pac-Islander' 'Other' 'Amer-Indian-Eskimo']  
  
sex : ['Female' 'Male']  
  
capital.gain : [ 0 99999 41310 34095 27828 25236 25124 22040 20051 18481 15831 15024  
15020 14344 14084 13550 11678 10605 10566 10520 9562 9386 8614 7978  
7896 7688 7443 7430 7298 6849 6767 6723 6514 6497 6418 6360  
6097 5721 5556 5455 5178 5060 5013 4934 4931 4865 4787 4687  
4650 4508 4416 4386 4101 4064 3942 3908 3887 3818 3781 3674  
3471 3464 3456 3432 3418 3411 3325 3273 3137 3103 2993 2977  
2964 2961 2936 2907 2885 2829 2653 2635 2597 2580 2538 2463  
2414 2407 2387 2354 2346 2329 2290 2228 2202 2176 2174 2105  
2062 2050 2036 2009 1848 1831 1797 1639 1506 1471 1455 1424  
1409 1173 1151 1111 1086 1055 991 914 594 401 114]  
  
capital.loss : [4356 3900 3770 3683 3004 2824 2754 2603 2559 2547 2489 2472 2467 2457  
2444 2415 2392 2377 2352 2339 2282 2267 2258 2246 2238 2231 2206 2205  
2201 2179 2174 2163 2149 2129 2080 2057 2051 2042 2002 2001 1980 1977  
1974 1944 1902 1887 1876 1848 1844 1825 1816 1762 1755 1741 1740 1735  
1726 1721 1719 1672 1669 1668 1651 1648 1628 1617 1602 1594 1590 1579  
1573 1564 1539 1504 1485 1411 1408 1380 1340 1258 1138 1092 974 880  
810 653 625 419 323 213 155 0]  
  
hours.per.week : [40 18 45 20 60 35 55 76 50 42 25 32 90 48 15 70 52 72 39 6 65 12 80 67  
99 30 75 26 36 10 84 38 62 44 8 28 59 5 24 57 34 37 46 56 41 98 43 63  
1 47 68 54 2 16 9 3 4 33 23 22 64 51 19 58 53 96 66 21 7 13 27 11  
14 77 31 78 49 17 85 87 88 73 89 97 94 29 82 86 91 81 92 61 74 95]  
  
native.country : ['United-States' '?' 'Mexico' 'Greece' 'Vietnam' 'China' 'Taiwan' 'India'  
'Philippines' 'Trinidad&Tobago' 'Canada' 'South' 'Holand-Netherlands'  
'Puerto-Rico' 'Poland' 'Iran' 'England' 'Germany' 'Italy' 'Japan' 'Hong'  
'Honduras' 'Cuba' 'Ireland' 'Cambodia' 'Peru' 'Nicaragua'  
'Dominican-Republic' 'Haiti' 'El-Salvador' 'Hungary' 'Columbia'  
'Guatemala' 'Jamaica' 'Ecuador' 'France' 'Yugoslavia' 'Scotland'  
'Portugal' 'Laos' 'Thailand' 'Outlying-US(Guam-USVI-etc)']  
  
income : ['<=50K' '>50K']
```

In [9]:

```
pd.isnull(df).sum()
```

Check for Null Data

Out[9]:

```
age          0
workclass    0
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   0
relationship 0
race         0
sex          0
capital.gain  0
capital.loss  0
hours.per.week 0
native.country 0
income       0
dtype: int64
```

numerical attributes

In [10]:

```
num_attributes = df.select_dtypes(include=['int'])
print(num_attributes.columns)
```

Identify Numeric features

```
Index(['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss',
      'hours.per.week'],
      dtype='object')
```

In [11]:

```
num_attributes.describe()
```

describe about the numerical columns(like mean,

Out[11]:

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

categorical_attributes

In [12]:

```
categorical_attributes = df.select_dtypes(include=['object'])
print(categorical_attributes.columns)
```

Identify Categorical

```
Index(['workclass', 'education', 'marital.status', 'occupation',
      'relationship', 'race', 'sex', 'native.country', 'income'],
      dtype='object')
```

In [13]:

```
categorical_attributes.describe()
```

Out[13]:

	workclass	education	marital.status	occupation	relationship	race	sex	native.country	income
count	32561	32561	32561	32561	32561	32561	32561	32561	32561
unique	9	16	7	15	6	5	2	42	2
top	Private	HS-grad	Married-civ-spouse	Prof-specialty	Husband	White	Male	United-States	<=50K
freq	22696	10501	14976	4140	13193	27816	21790	29170	24720

encoded number for each field

In [14]:

```
df.workclass.value_counts()
```

Out[14]:

```
Private          22696
Self-emp-not-inc 2541
Local-gov        2093
?                1836
State-gov        1298
Self-emp-inc     1116
Federal-gov      960
Without-pay      14
Never-worked      7
Name: workclass, dtype: int64
```

In [15]:

```
df.groupby('education')['education.num'].unique().sort_values()
# Education.num shows the no from 1 to 16 on the education basis, We are getting the same info from the education
```

Out[15]:

```
education
Preschool      [1]
1st-4th        [2]
5th-6th        [3]
7th-8th        [4]
9th            [5]
10th           [6]
11th           [7]
12th           [8]
HS-grad        [9]
Some-college   [10]
Assoc-voc      [11]
Assoc-acdm     [12]
Bachelors      [13]
Masters        [14]
Prof-school    [15]
Doctorate      [16]
Name: education.num, dtype: object
```

In [16]:

```
df['marital.status'].value_counts()
```

Out[16]:

```
Married-civ-spouse      14976
Never-married           10683
Divorced                 4443
Separated               1025
Widowed                 993
Married-spouse-absent   418
Married-AF-spouse       23
Name: marital.status, dtype: int64
```

In [17]:

```
df.occupation.value_counts()
```

Out[17]:

```
Prof-specialty      4140
Craft-repair        4099
Exec-managerial     4066
Adm-clerical        3770
Sales               3650
Other-service       3295
Machine-op-inspct   2002
?                  1843
Transport-moving    1597
Handlers-cleaners   1370
Farming-fishing     994
Tech-support        928
Protective-serv     649
Priv-house-serv     149
Armed-Forces        9
Name: occupation, dtype: int64
```

In [18]:

```
df.relationship.value_counts()
```

Out[18]:

```
Husband      13193
Not-in-family 8305
Own-child    5068
Unmarried    3446
Wife         1568
Other-relative 981
Name: relationship, dtype: int64
```

In [19]:

```
df.race.value_counts()
```

Out[19]:

```
White      27816
Black      3124
Asian-Pac-Islander 1039
Amer-Indian-Eskimo 311
Other       271
Name: race, dtype: int64
```

In [20]:

```
df.sex.value_counts()
```

Out[20]:

```
Male      21790
Female    10771
Name: sex, dtype: int64
```

In [21]:

```
df['native.country'].value_counts()
```

Out[21]:

```
United-States      29170
Mexico              643
?                  583
Philippines        198
Germany            137
Canada             121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba               95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Guatemala          64
Japan              62
Poland             60
Columbia           59
Taiwan             51
Haiti              44
Iran               43
Portugal           37
Nicaragua          34
Peru               31
Greece            29
France            29
Ecuador           28
Ireland           24
Hong              20
Cambodia          19
Trinidad&Tobago    19
Laos              18
Thailand           18
Yugoslavia        16
Outlying-US(Guam-USVI-etc) 14
Hungary           13
Honduras          13
Scotland          12
Holand-Netherlands 1
Name: native.country, dtype: int64
```

In [22]:

```
df.income.value_counts()
```

Out[22]:

```
<=50K    24720
>50K      7841
Name: income, dtype: int64
```


In [23]:

```
df.isin(['?']).sum(axis=0)
```

finding out total number of "?" sym

Out[23]:

```
age          0
workclass    1836
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   1843
relationship 0
race         0
sex          0
capital.gain 0
capital.loss 0
hours.per.week 0
native.country 583
income       0
dtype: int64
```

In [24]:

```
df.groupby(['income', 'occupation']).size()
```

Out[24]:

```
income  occupation
<=50K   ?          1652
        Adm-clerical 3263
        Armed-Forces 8
        Craft-repair 3170
        Exec-managerial 2098
        Farming-fishing 879
        Handlers-cleaners 1284
        Machine-op-inspct 1752
        Other-service 3158
        Priv-house-serv 148
        Prof-specialty 2281
        Protective-serv 438
        Sales 2667
        Tech-support 645
        Transport-moving 1277
>50K    ?          191
        Adm-clerical 507
        Armed-Forces 1
        Craft-repair 929
        Exec-managerial 1968
        Farming-fishing 115
        Handlers-cleaners 86
        Machine-op-inspct 250
        Other-service 137
        Priv-house-serv 1
        Prof-specialty 1859
        Protective-serv 211
        Sales 983
        Tech-support 283
        Transport-moving 320
dtype: int64
```

Reformatting Column

converting dependent col into categor- assuming <=50K - 0, >50K - 1

In [25]:

```
df['income'] = df['income'].map({'<=50K':0, '>50K':1})
#df['income'].replace({'<=50K':0, '>50K':1}, inplace=True)
df.head(10)
```

Out[25]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.g
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	
7	74	State-gov	88638	Doctorate	16	Never-married	Prof-specialty	Other-relative	White	Female	
8	68	Federal-gov	422013	HS-grad	9	Divorced	Prof-specialty	Not-in-family	White	Female	
9	41	Private	70037	Some-college	10	Never-married	Craft-repair	Unmarried	White	Male	

Event rate

In [26]:

```
df['income'].value_counts(normalize=True)*100
```

Out[26]:

```
0    75.919044
1    24.080956
Name: income, dtype: float64
```

In [27]:

```
df.groupby(['income', 'sex']).size()
```

Out[27]:

```
income  sex
0       Female    9592
        Male     15128
1       Female    1179
        Male     6662
dtype: int64
```

Outliers detection

Q1 Q3 IQR I R IIR

In [28]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num        32561 non-null  int64
5   marital.status       32561 non-null  object
6   occupation           32561 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                 32561 non-null  object
10  capital.gain         32561 non-null  int64
11  capital.loss         32561 non-null  int64
12  hours.per.week       32561 non-null  int64
13  native.country       32561 non-null  object
14  income              32561 non-null  int64
dtypes: int64(7), object(8)
memory usage: 3.7+ MB
```

In [29]:

```
col_for_outliers=['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week']
```

In [30]:

```
summary_pre_outliers_detection = df[col_for_outliers].describe()
```

In [31]:

```
summary_pre_outliers_detection
```

Out[31]:

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

In [32]:

```
for i in col_for_outliers:
    Q1 = np.percentile(df[i], 25)
    Q3 = np.percentile(df[i], 75)
```

In [33]:

```
IQR = Q3-Q1
```

In [34]:

```
LB = Q1-1.5*IQR
```

In [35]:

```
UB = Q3+1.5*IQR
```

In [36]:

```
df[i] = np.where(df[i] < LB, LB, df[i])
```

In [37]:

```
df[i] = np.where(df[i] > UB, UB, df[i])
```

In [38]:

```
summary_post_outliers_detection = df[col_for_outliers].describe()
```

In [39]:

```
summary_post_outliers_detection
```

Out[39]:

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	41.202451
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	6.187005
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	32.500000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	52.500000

In [40]:

```
df.head(10)
```

Out[40]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.g
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	
7	74	State-gov	88638	Doctorate	16	Never-married	Prof-specialty	Other-relative	White	Female	
8	68	Federal-gov	422013	HS-grad	9	Divorced	Prof-specialty	Not-in-family	White	Female	
9	41	Private	70037	Some-college	10	Never-married	Craft-repair	Unmarried	White	Male	

Univariate Analysis

For col workclass and occupation we are replacing ? with never-worked

In [41]:

```
#print(df.replace("?",np.nan, inplace = True)) # this will replace the value which is on the
```

In [42]:

```
##(Univariate analysis)
df['workclass']=df['workclass'].replace(to_replace="?",value="never-worked")
df.workclass.value_counts() # this will replace the '?' from workclass to Unemployed
```

Out[42]:

```
Private          22696
Self-emp-not-inc  2541
Local-gov        2093
never-worked     1836
State-gov        1298
Self-emp-inc     1116
Federal-gov       960
Without-pay       14
Never-worked       7
Name: workclass, dtype: int64
```

In [43]:

```
#df['workclass'].replace(to_replace = ['?', 'Self-emp-not-inc', 'Without-pay', 'Never-worked'], value = 'no-income', inplace = True)
#df['workclass'].replace(to_replace = ['Local-gov', 'State-gov', 'Federal-gov'], value = 'gov', inplace = True)
#df['workclass'].replace(to_replace = 'Self-emp-inc', value = 'Self', inplace = True)
```

In [44]:

```
df.head()
```

Out[44]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.g
0	90	never-worked	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
2	66	never-worked	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	

In [45]:

```
df['occupation']=df['occupation'].replace(to_replace="?",  
                                           value="never-worked")  
#df['occupation'].replace(to_replace = ['?', 'Other-service'], value = 'Other', inplace = True)  
df.occupation.value_counts()
```

Out[45]:

```
Prof-specialty      4140  
Craft-repair        4099  
Exec-managerial     4066  
Adm-clerical        3770  
Sales               3650  
Other-service       3295  
Machine-op-inspct   2002  
never-worked        1843  
Transport-moving    1597  
Handlers-cleaners   1370  
Farming-fishing     994  
Tech-support        928  
Protective-serv     649  
Priv-house-serv     149  
Armed-Forces         9  
Name: occupation, dtype: int64
```

In [46]:

```
df.head()
```

Out[46]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.g
0	90	never-worked	77053	HS-grad	9	Widowed	never-worked	Not-in-family	White	Female	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
2	66	never-worked	186061	Some-college	10	Widowed	never-worked	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	

In [47]:

```
df['native.country'].value_counts()
```

Out[47]:

United-States	29170
Mexico	643
?	583
Philippines	198
Germany	137
Canada	121
Puerto-Rico	114
El-Salvador	106
India	100
Cuba	95
England	90
Jamaica	81
South	80
China	75
Italy	73
Dominican-Republic	70
Vietnam	67
Guatemala	64
Japan	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
Greece	29
France	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Hungary	13
Honduras	13
Scotland	12
Holand-Netherlands	1

Name: native.country, dtype: int64

In native.country we are replacing ? with mode

In [48]:

```
temp_mode = df['native.country'].mode()[0]  
df['native.country'] = df['native.country'].replace('?', temp_mode)
```

**United states appearing maximum no of times(29k), we can segregate the column in 2 parts
United states & others take other countries as one entity.**

In [49]:

```
# Loop for segregate the native.country into US and Other
col = df['native.country']

for i in col:
    print(i)
    if i != 'United-States':
        df['native.country'] = df['native.country'].replace({i: 'others'})
```

[illegible]

In [50]:

```
df['native.country'].value_counts()
```

Out[50]:

```
United-States    29753
others           2808
Name: native.country, dtype: int64
```

In [51]:

```
df.head()
```

Out[51]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain
0	90	never-worked	77053	HS-grad	9	Widowed	never-worked	Not-in-family	White	Female	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
2	66	never-worked	186061	Some-college	10	Widowed	never-worked	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	

In [52]:

```
#pd.crosstab(raw['native.country'], raw['income'],
#values=raw['hours.per.week'], aggfunc=np.mean)
```

In [53]:

```
df.isnull().sum()
```

Out[53]:

```
age                0
workclass          0
fnlwgt            0
education         0
education.num     0
marital.status    0
occupation        0
relationship      0
race              0
sex               0
capital.gain      0
capital.loss      0
hours.per.week    0
native.country    0
income            0
dtype: int64
```

In [54]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  float64
13  native.country        32561 non-null  object
14  income                32561 non-null  int64
dtypes: float64(1), int64(6), object(8)
memory usage: 3.7+ MB
```

In [55]:

```
df.isnull().sum()
```

Out[55]:

```
age          0
workclass    0
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   0
relationship 0
race         0
sex          0
capital.gain  0
capital.loss  0
hours.per.week 0
native.country 0
income       0
dtype: int64
```

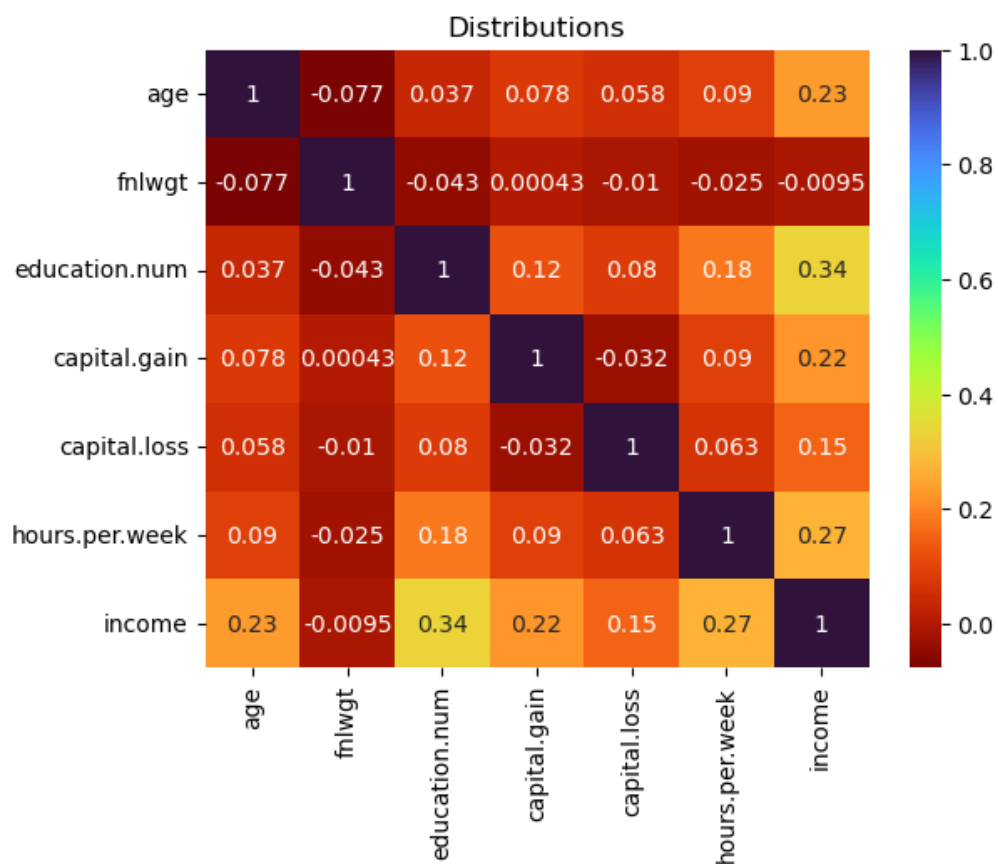
Bivariate analysis for continuous variables

In [56]:

```
corr_df = df.corr()
sns.heatmap(corr_df, xticklabels=corr_df, yticklabels=corr_df, cmap='turbo_r', annot=True)
plt.title('Distributions')
```

Out[56]:

```
Text(0.5, 1.0, 'Distributions')
```



In [57]:


```
# Creating a countplot of income across Marital Status
plt.figure(figsize=(15,8))
graph1=sns.countplot(x='marital.status',hue='income',data=df)

for i in graph1.containers:
    graph1.bar_label(i)
plt.legend(loc='upper right')
plt.title("Income across Marital Status", fontdict={'fontsize': 20, 'fontweight': 'bold'})

#Based on the graph analysis it is clear that Never-married peoples salary is less than 50K but more in total
#and same for marries-civ-spouse
```

Out[57]:

Text(0.5, 1.0, 'Income across Marital Status')

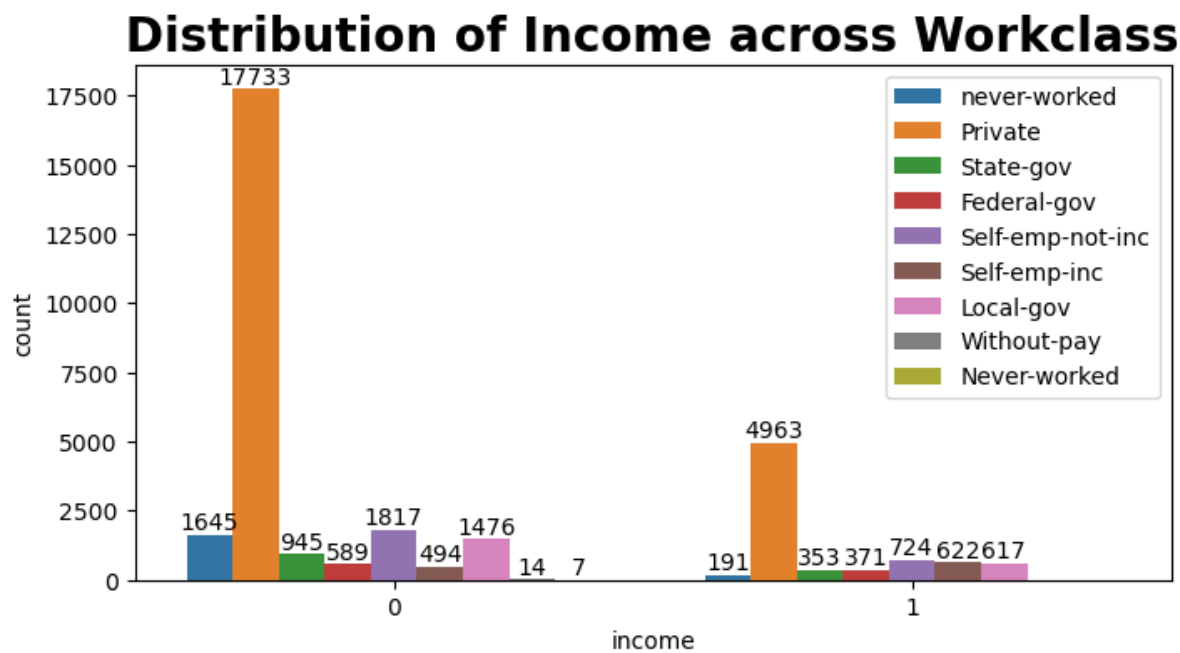


In [58]:

```
# Creating a countplot of income across workclass
plt.figure(figsize=(8,4))
graph3=sns.countplot(x='income',hue='workclass',data=df)
for i in graph3.containers:
    graph3.bar_label(i)
plt.legend(loc='upper right')
plt.title('Distribution of Income across Workclass', fontdict={'fontsize': 20, 'fontweight': 'bold'})
#From the 'Workclass' countplot graph it can be analyse that most of the people are working privately with
#less than 50K income.
```

Out[58]:

Text(0.5, 1.0, 'Distribution of Income across Workclass')

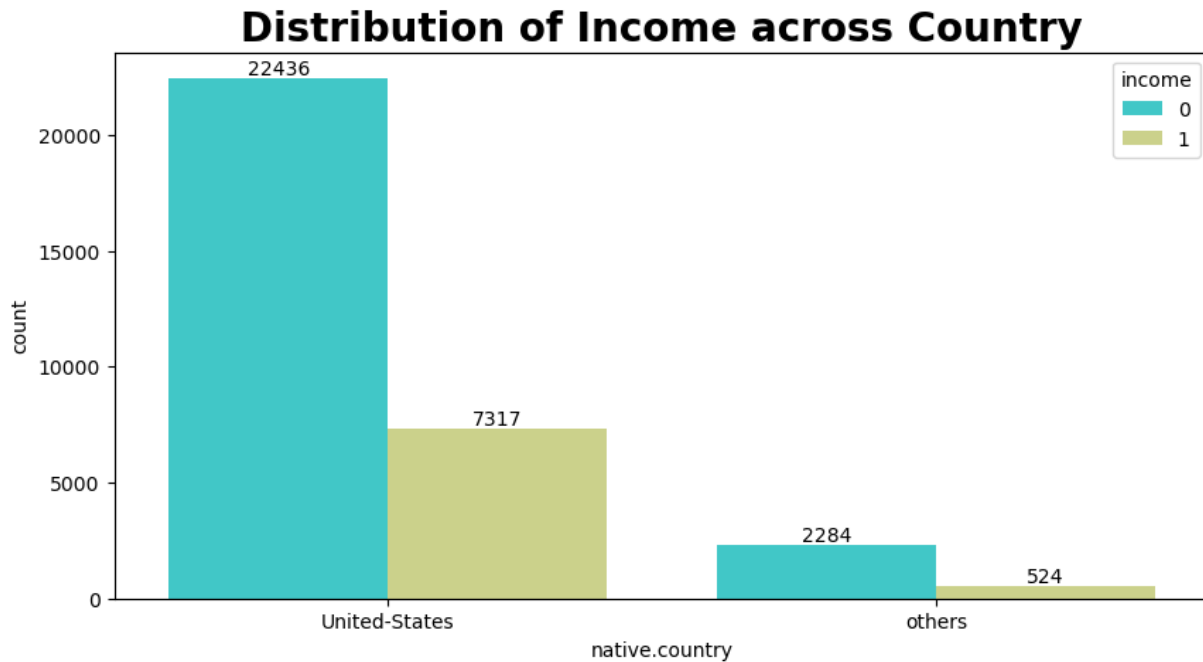


In [59]:

```
# Creating a countplot of income across country
plt.figure(figsize=(10,5))
graph4=sns.countplot(x="native.country", hue="income", data=df,palette='rainbow')
for i in graph4.containers:
    graph4.bar_label(i)
plt.title('Distribution of Income across Country', fontdict={'fontsize': 20, 'fontweight': 'bold'})
#From the 'native country' countplot graph it can be analyse that maximum number of people working in united s
#and had a income <=50K
```

Out[59]:

Text(0.5, 1.0, 'Distribution of Income across Country')




In [60]:

```
# Creating a countplot of income across occupation
plt.figure(figsize=(25,4))
graph=sns.countplot(data=df, x='occupation', hue='income')
for i in graph.containers:
    graph.bar_label(i)
plt.title('Distribution of Income across Occupation', fontdict={'fontsize': 20, 'fontweight': 'bold'})
#From the 'Occupation' countplot graph it can be analyse that most of the peoples work as Adm-Clerical, Craft
```


Out[60]:

Text(0.5, 1.0, 'Distribution of Income across Occupation')



In [61]:

```
# Creating a countplot for 'Hours per week'
plt.figure(figsize=(25,6))
graph=sns.countplot(x="hours.per.week",data=df)
for i in graph.containers:
    graph.bar_label(i)
plt.title('Distribution of Hours of work per week', fontdict={'fontsize': 20, 'fontweight': 'bold'})
# it is graph from the graph that majority of the people are working 40 hours in a week.
```




In [62]:

```
# Creating a boxplot for workclass
plt.figure(figsize=(14,4))
sns.boxplot(x="workclass", y="fnlwgt", data=df,palette='rainbow')
#Outliers present in all the workclass w.r.t final weight
```

Out[62]:

<AxesSubplot:xlabel='workclass', ylabel='fnlwgt'>



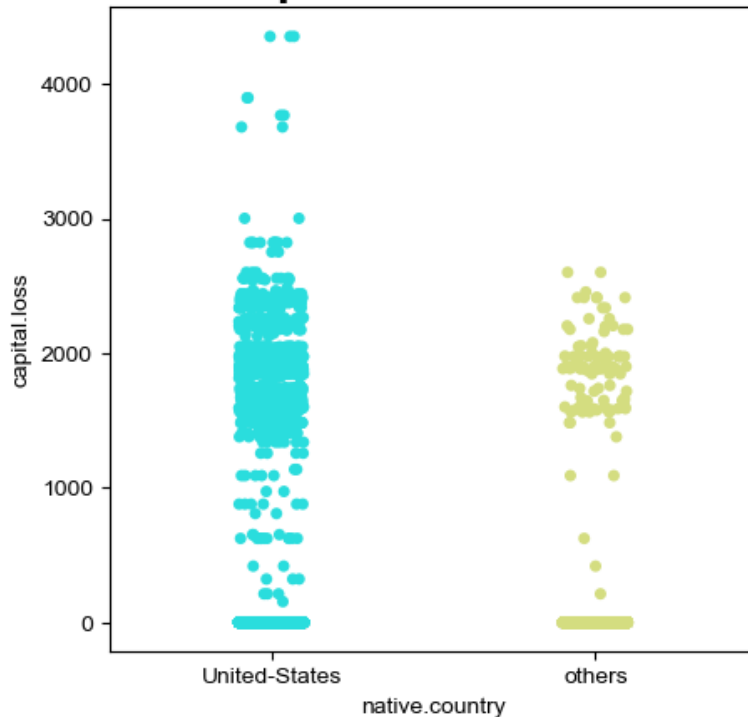
In [63]:

```
# Creating a strip plot of native country accross capital loss
plt.figure(figsize=(5,5))
sns.stripplot(x="native.country", y="capital.loss", data=df,palette='rainbow')
sns.set(rc={'figure.figsize':(30,20)})
plt.title('Distribution of capital loss accross native country', fontdict={'fontsize': 20, 'fontweight': 'bold'})
#capital loss is highest in united states
```

Out[63]:

Text(0.5, 1.0, 'Distribution of capital loss accross native country')

Distribution of capital loss accross native country

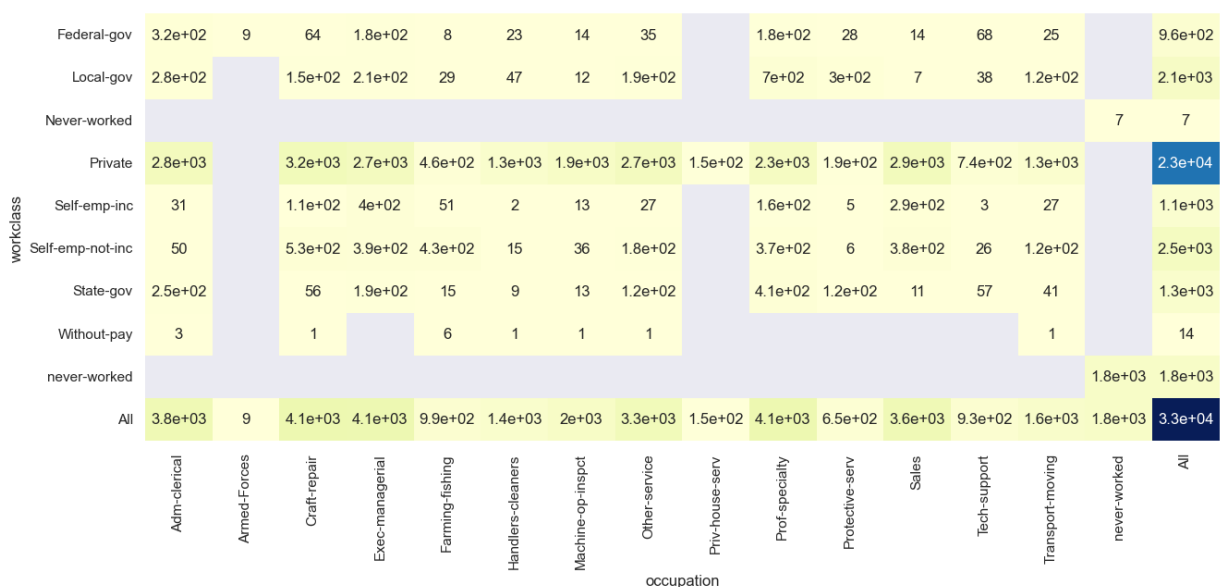


In [64]:

```
# Creating a heat map of occupation accross workclass
plt.figure(figsize=(15,6))
sns.heatmap(pd.crosstab(df.workclass, df.occupation, margins=True, values=df.income, aggfunc=pd.Series.count),
            annot=True, cbar=False)
```


Out[64]:

<AxesSubplot: xlabel='occupation', ylabel='workclass'>



In [65]:

```
# Creating a distribution plot for 'Age'
age = df['age'].value_counts()
plt.figure(figsize=(10, 5))
sns.distplot(df['age'], bins=20)
plt.title('Distribution of Age', fontdict={'fontsize': 20, 'fontweight': 'bold'})
plt.xlabel('Age')
plt.ylabel('Number of people')
plt.show()
#From the graph shown below it can be predicted that people of age 25-45 are more in population.
```




In [66]:

```
# Creating a pie chart for 'Race'
plt.figure(figsize=(8,8))
labels = df['race'].value_counts().index
values = df['race'].value_counts().values
colors = df['race']
plt.pie(values, labels=labels, autopct="%1.2f%%")
plt.title('Race', fontdict={'fontsize': 20, 'fontweight': 'bold'})
plt.legend()
#based on the graph analysis it is clear that 85.43% people are white
```


Out[66]:

<matplotlib.legend.Legend at 0x172fe911ca0>




In [67]:

```
# Creating a pie chart for 'Marital status'
marital = df['marital.status'].value_counts()
plt.figure(figsize=(10, 7))
plt.pie(marital.values, labels=marital.index, startangle=10, explode=(0, 0.20, 0, 0, 0, 0, 0), shadow=True, autopct='%1.1f%%')
plt.title('Marital distribution', fontdict={'fontname': 'Monospace', 'fontsize': 20, 'fontweight': 'bold'})
plt.legend()
plt.axis('equal')
plt.show()
#Based on the graph analysis it can be conclude that Mostly population are Married-civ-spouse
```




In [68]:

```
# Creating a pie chart for 'Gender'
label=df.sex.value_counts().index
count=df.sex.value_counts().values
plt.figure(1, figsize=(4,4))
plt.pie(count,labels=label,autopct='%1.1f%%')
plt.title('Gender', fontdict={'fontsize': 20, 'fontweight': 'bold'})
plt.show()
#Based on the graph analysis it is clear that males percentage is more then females in short males are more t
```



In [69]:

```
# Creating a donut chart for 'Relationship'
relation = df['relationship'].value_counts()
plt.figure(figsize=(16, 10))
plt.pie(relation.values, labels=relation.index, startangle=50, autopct='%1.1f%%')
centre_circle = plt.Circle((0, 0), 0.7, fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)
plt.title('Relationship distribution', fontdict={'fontsize': 20, 'fontweight': 'bold'})
plt.axis('equal')
plt.legend()
plt.show()
#From the 'Relationship Distribution graph' it can be analyse that most of the husbands are working.
```



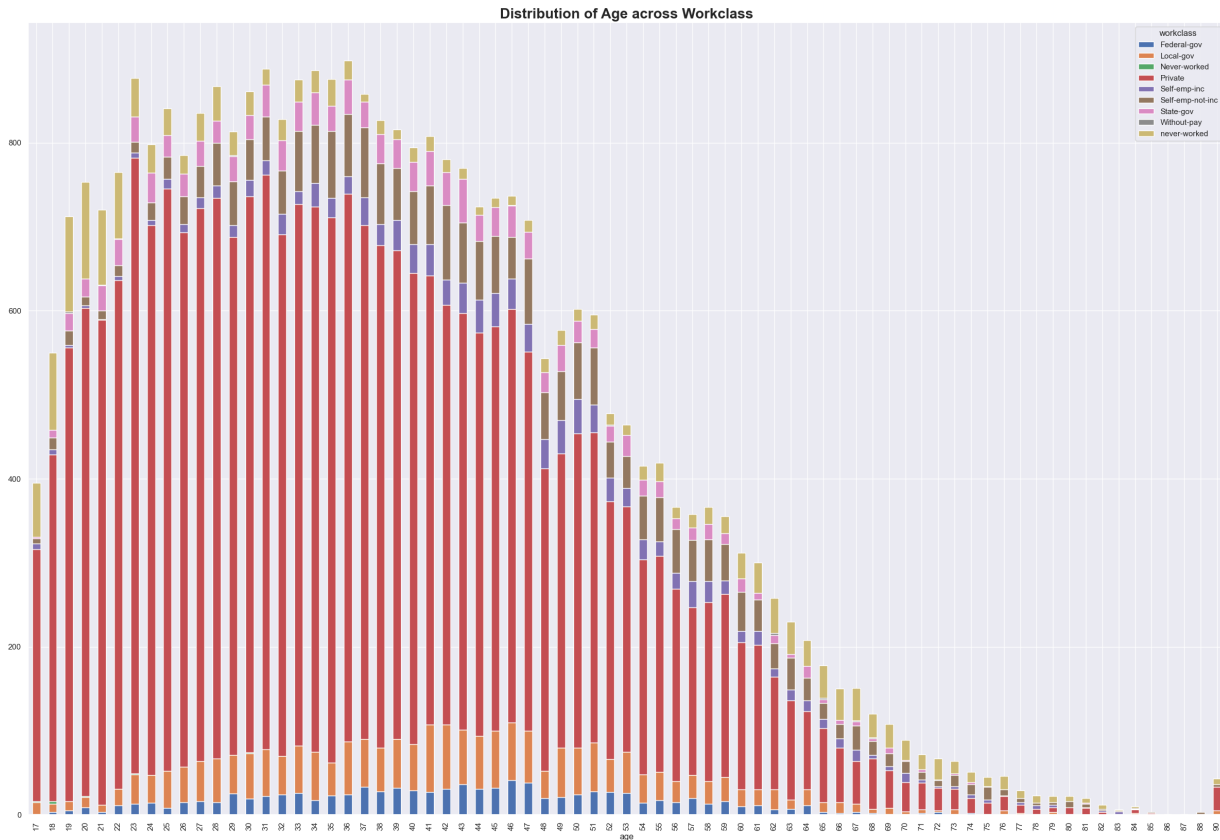
In [70]:

```
# Creating a barplot of Age across Workclass
plt.figure(figsize=(25,18))
df.groupby(['age', 'workclass']).size().unstack().plot(kind='bar', stacked=True)
plt.title('Distribution of Age across Workclass', fontdict={'fontsize': 20, 'fontweight': 'bold'})
#based on the graph analysis it is clear that most of the people are working privately in each and every age.
```

Out[70]:

Text(0.5, 1.0, 'Distribution of Age across Workclass')

<Figure size 2500x1800 with 0 Axes>



In [71]:

```
#from matplotlib.pyplot import figure

#category_var=['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week','workclass',
# 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']

#for i in category_var:
#    #figure()
#    #sns.barplot(y=df['income_num'], x=df[i])
#    #sns.set(rc={'figure.figsize':(25,10)})
```

In [72]:

```
from matplotlib.pyplot import figure
#Creating a count plot for the following fields.
category_var=['workclass', 'education', 'occupation', 'relationship', 'race', 'sex', 'native.country',
              'hours.per.week', 'marital.status']
for i in category_var:
    figure()
    graph=sns.countplot(data=df, x=df[i], hue='income')
    for a in graph.containers:
        graph.bar_label(a)
    #sns.barplot(y=df['income_num'], x=df[i])
    sns.set(rc={'figure.figsize':(25,10)})
#GRAPH 1 --- From the 'Workclass' countplot graph it can be analyse that most of the people are working private and had a salary less than 50K.
#Less than 50K income.

#GRAPH 2 --- From the 'Education' countplot graph it can be analyse that most of the peoples are High School or below and had a salary less than or equal to 50K.
#whose income is less than or equal to 50K.

#GRAPH 3 --- From the 'Occupation' countplot graph it can be analyse that most of the peoples work as Adm-Clerical and had a salary less than or equal to 50K.
#Craft repair.

#GRAPH 4 --- From the 'Relationship' countplot graph it can be analyse that maximum number of people does not have a family and had a salary less than or equal to 50K.
#with their family and had a salary less than or equal to 50K.

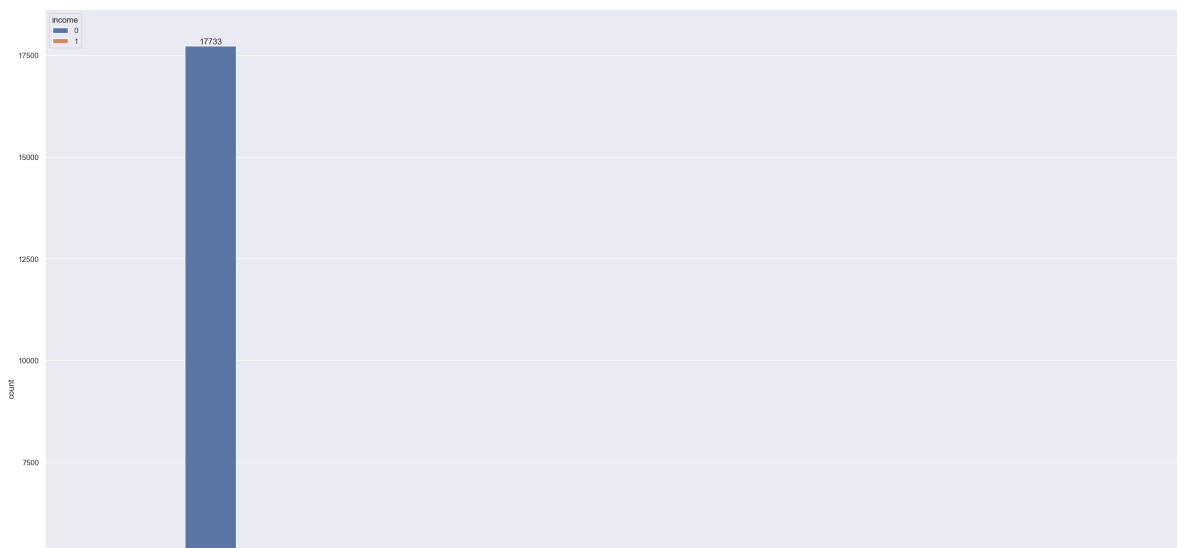
#GRAPH 5 --- From the 'Race' countplot graph it can be analyse that most of the people are white and have a income less than or equal to 50K.
#and had a income <=50K

#GRAPH 6 --- From the 'Sex' countplot graph it can be analyse that maximum population who is working is male and had a salary less than or equal to 50K.
#<=50K

#Graph 7 --- From the 'native country' countplot graph it can be analyse that maximum number of people working is native born and had a income <=50K
#and had a income <=50K


#Graph 8 --- From the 'hours per week' countplot graph it can be analyse that people mostly work 40 hours a week and had a salary less than or equal to 50K.
#salary <=50K

#Graph 9 --- From the 'Never married' countplot graph it can be analyse that people who are unmarried are working and had a salary less than or equal to 50K.
# rather tan others.
```



In [73]:

```
#Creating a plot for the following fields.
g = sns.FacetGrid(df, col='income')
g = g.map(sns.distplot, "age")
plt.show()
#From the graph shown below it can be analyse that people belonging to age 20-40 are more and have there income
#From the graph shown below it can be analyse that people belonging to age 30-50 are more and have there income
```




In [74]:

```
plt.figure(figsize=(20,4))
graph=sns.countplot(data=df, x='hours.per.week', hue='sex')
for i in graph.containers:
    graph.bar_label(i)
plt.title('Distribution of gender across working hours', fontdict={'fontsize': 20, 'fontweight': 'bold'})
```

Out[74]:

Text(0.5, 1.0, 'Distribution of gender across working hours')




In [75]:

```
plt.figure(figsize=(25,4))
graph=sns.countplot(data=df, x='hours.per.week', hue='income')
for i in graph.containers:
    graph.bar_label(i)
plt.title('Distribution of Income across working hours', fontdict={'fontsize': 20, 'fontweight': 'bold'})
```

Out[75]:

Text(0.5, 1.0, 'Distribution of Income across working hours')

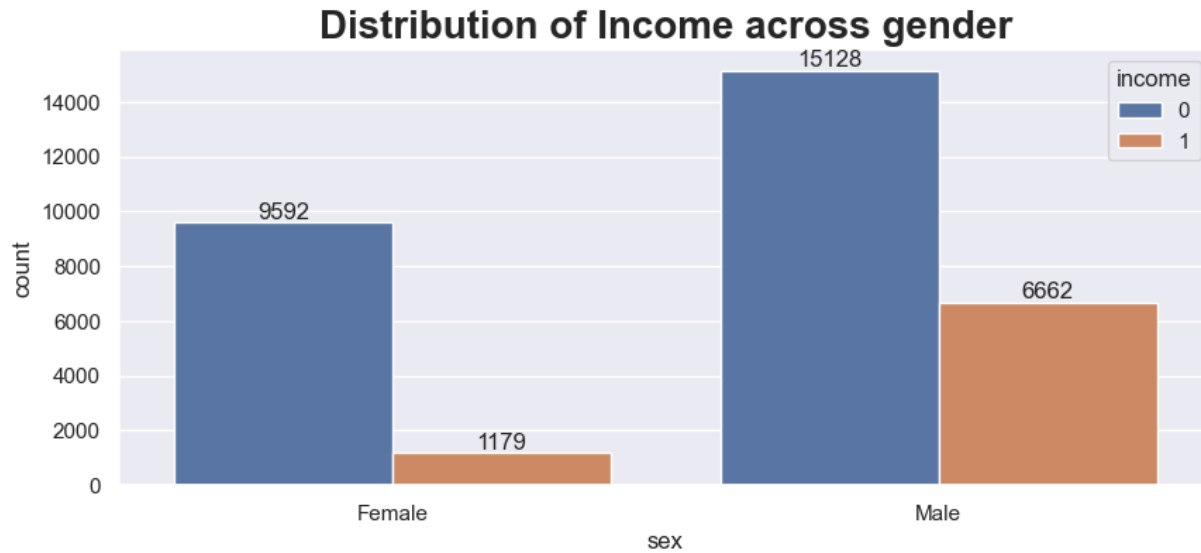


In [76]:

```
plt.figure(figsize=(10,4))
graph=sns.countplot(data=df, x='sex', hue='income')
for i in graph.containers:
    graph.bar_label(i)
plt.title('Distribution of Income across gender', fontdict={'fontsize': 20, 'fontweight': 'bold'})
```


Out[76]:

Text(0.5, 1.0, 'Distribution of Income across gender')



In [77]:


```
# Creating a barplot for 'Income'
income = df['income'].value_counts()
plt.figure(figsize=(7, 5))
graph=sns.barplot(income.index, income.values)
for a in graph.containers:
    graph.bar_label(a)
plt.title('Distribution of Income', fontdict={'fontsize': 20, 'fontweight': 'bold'})
plt.xlabel('Income')
plt.ylabel('Number of people')
plt.show()
#From the graph shown below it is clear that maximum number of people's salary is less than or equal to 50 thousand
```



Statistical Tests

In [78]:

```
#Boxplot analysis between age and income
fig = plt.figure(figsize=(10,10))
sns.boxplot(x="income", y="age", data=df)
plt.show()
#Outliers present in both the income group(<=50k and >50k) wrt "age" attribute.
#Income group(<=50k) has lower median "age"(34 year) than the Income group(>50k) which has median "age"(43 year)
#For Income group(<=50k) , Interquartile range(IQR) is between [25,46] (Long range)
#For Income group(>50k) , Interquartile range(IQR) is between [35,50] (shorter range)
```



In [79]:

```
from scipy import stats
import random
from scipy.stats import ttest_ind, ttest_rel
#Hypothesis test (to test the relationship between 'income' & 'age' )
df = df[(np.abs(stats.zscore(df["age"]))) < 3]
income_1 = df[df['income']==1]['age']
income_0 = df[df['income']==0]['age']

income_0 = income_0.values.tolist()
income_0 = random.sample(income_0, 100)
income_1 = income_1.values.tolist()
income_1 = random.sample(income_1, 100)

ttest,pval = ttest_ind(income_1,income_0,equal_var = False)
print("ttest",ttest)
print('p value',pval)

if pval <0.05:
    print("we reject null hypothesis")
else:
    print("we accept null hypothesis")


#Using statistical analysis, we conclude that there is a significant difference in the mean ages of income groups
#and income group <=50k.It means that age has some contribution to the distinguish income groups.
```

ttest 4.611001896522119
p value 7.344483569695718e-06
we reject null hypothesis

In [80]:

```
#Boxplot relation between 'income' and 'hours.per.week'
fig = plt.figure(figsize=(10,10))
sns.boxplot(x="income", y="hours.per.week", data=df)
plt.show()
```

```
#The median "hours.per.week" for income group who earns >50k is greater than the income group who earns <=50k
#The boxplot for Income group who earns <=50k has small range ~[28,48].
#The boxplot for Income group who earns >50k has Large range ~[25,65].
#Income group who earns >50k have flexible working hours
#More Outliers present in the Income group who earns <=50k.
```



In [81]:

```
#Hypothesis test (to test the relationship between 'income' & 'hours.per.week' )
df = df[(np.abs(stats.zscore(df["hours.per.week"]))) < 3]
income_1 = df[df['income']==1]["hours.per.week"]
income_0 = df[df['income']==0]["hours.per.week"]

income_0 = income_0.values.tolist()
income_0 = random.sample(income_0, 100)
income_1 = income_1.values.tolist()
income_1 = random.sample(income_1, 100)

ttest,pval = ttest_ind(income_1,income_0,equal_var = False)
print("ttest",ttest)
print('p value',format(pval, '.70f'))

if pval <0.05:
    print("we reject null hypothesis")
else:
    print("we accept null hypothesis")

#We can conclude that there is difference in Mean of income group >50k and income group <=50k.
#It means that hours-per-week has some contribution to the distinguish income groups.
```

```
ttest 2.9916437116170367
p value 0.0031279895728443369261329021213668966083787381649017333984375000000000
we reject null hypothesis
```

In [82]:

```
from scipy import stats
```

In [83]:

```
df.head(8)
```

Out[83]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.g
2	66	never-worked	186061	Some-college	10	Widowed	never-worked	Unmarried	Black	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	
7	74	State-gov	88638	Doctorate	16	Never-married	Prof-specialty	Other-relative	White	Female	
8	68	Federal-gov	422013	HS-grad	9	Divorced	Prof-specialty	Not-in-family	White	Female	
9	41	Private	70037	Some-college	10	Never-married	Craft-repair	Unmarried	White	Male	

In [84]:

```
df.tail(8)
```

Out[84]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	ca
32553	43	Private	84661	Assoc-voc	11	Married-civ-spouse	Sales	Husband	White	Male	
32554	32	Private	116138	Masters	14	Never-married	Tech-support	Not-in-family	Asian-Pac-Islander	Male	
32555	53	Private	321865	Masters	14	Married-civ-spouse	Exec-managerial	Husband	White	Male	
32556	22	Private	310152	Some-college	10	Never-married	Protective-serv	Not-in-family	White	Male	
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	
32558	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
32559	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	
32560	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	

In [85]:

```
Sales = df[(df['occupation'] == 'Sales')]
Sales.shape
```

Out[85]:

```
(3639, 15)
```

In [86]:

```
Adm=df[(df['occupation'] == 'Adm-clerical')]
Adm.shape
```

Out[86]:

```
(3759, 15)
```

In [87]:

```
Sales['income']=Sales['income'].sample(28)
Adm['income']=Adm['income'].sample(28)
```

In [88]:

```
print(np.mean(Sales['age']))
print(np.mean(Adm['age']))
```

```
37.21077219016213
36.82255919127427
```

In [89]:

```
tvalue,pvalue=stats.ttest_ind(Sales['age'], Adm['age'])
```

In [90]:

```
pvalue
```

Out[90]:

```
0.21773563264672965
```

In [91]:

```
tvalue
```

Out[91]:

```
1.2326765836747073
```

In [92]:

```
H0="Mean value of both distributions is same"  
H1="Mean value is different"
```

In [93]:

```
if pvalue>=0.05:  
    print(H0)  
else:  
    print(H1)
```

Mean value of both distributions is same

In [94]:

```
df['age'].head(20)
```

Out[94]:

```
2    66  
3    54  
4    41  
5    34  
6    38  
7    74  
8    68  
9    41  
10   45  
11   38  
12   52  
13   32  
14   51  
15   46  
16   45  
17   57  
18   22  
19   34  
20   37  
21   29
```

Name: age, dtype: int64

In [95]:

```
# ztest
```

In [96]:

```
capital_gain=df[df['capital.gain']==0]['income']  
capital_loss=df[df['capital.loss']>0]['income']
```

In [97]:

```
from statsmodels.stats.weightstats import ztest
```

In [98]:

```
z_score,p_val = ztest(capital_gain,capital_loss)
if p_val>0.05:
    print('Ho:hypothesis is true(there is no effect in income)')
else:
    print('H1:hypothesis in not true(there is effect on income)')
```

H1:hypothesis in not true(there is effect on income)

In [99]:

```
print(p_val)
```

8.991549662209643e-173

In [100]:

#Conclusions

#We did the entire EDA process for this dataset from looking at the head of the dataset to get the insights of every feature whether it is univariate analysis or the bivariate analysis and along with getting the insights numerically we also have used two one of the most interactive visualization libraries i.e. Count plot, Bar plot, heatmap, line plot, distplot, etc..

#75.92% of them are belong to income group 1 (who earns more than 50k) and 24.08% fall under the income group less than 50k).

#Females have more flexible working hours per week in the income groups who earns <=50k.

#Males have more flexible working hours per week in the income groups who earns >50k.

#Generally people can be seen working for 30 hours to 40 hours per week and they are not living with their families.

#For "female" earning more than 50k is rare with only 3.57% of all observations But for male, 19.99% of all people earn more than 50k .

#self-emp-inc workclass is only where more people earn >50k(belong to income group 1).

#People having degree doctorate,prof-school,masters are making salary more than 50K

#The people who are working mostly are unmarried probably belong to United states and working in private sector occupation is Adm-clerical.

#people of age group 25-45 are mostly working.

#most of the people who are working privately are high school graduate.

#maximum people race is white and males are more than female in whole population.

#Males are doing there jobs more than females and mostly males who are working are husbands.

In [101]:

```
from sklearn.preprocessing import LabelEncoder
```

In [102]:

```
for col in df.columns:
    if df[col].dtypes == 'object':
        encoder = LabelEncoder()
        df[col] = encoder.fit_transform(df[col])
```

Model Building

In [103]:

```
X=df.drop(['income'],axis=1)
Y=df['income']
```

Feature scaling

In [104]:

#As we have many features contains categorical variable so we are using pandas get_dummies function to convert

In [105]:

```
df= pd.get_dummies(df,drop_first=True)
pd.set_option('display.max_columns',100)#to display all columns
```

In [106]:

```
df.head(10)
#Now our data set has been transform into numeric.
```

Out[106]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain
2	66	8	186061	15	10	6	14	4	2	0	0
3	54	3	140359	5	4	0	6	4	4	0	0
4	41	3	264663	15	10	5	9	3	4	0	0
5	34	3	216864	11	9	0	7	4	4	0	0
6	38	3	150601	0	6	5	0	4	4	1	0
7	74	6	88638	10	16	4	9	2	4	0	0
8	68	0	422013	11	9	0	9	1	4	0	0
9	41	3	70037	15	10	4	2	4	4	1	0
10	45	3	172274	10	16	0	9	4	2	0	0
11	38	5	164526	14	15	4	9	1	4	1	0

In [107]:

```
df.shape
```

Out[107]:

(32440, 15)

In [108]:

*# Now our almost data values is 0 and 1 except few features Like “Age”,’Fnlwgt’,’Education_num’,’Hours_per_week’
we can use standard scaler we and convert those features in same scale.*

In [109]:

```
from sklearn.preprocessing import StandardScaler
```

In [110]:

```
scaler = StandardScaler()
train_col_sacle = df[['age', 'fnlwgt', 'education.num', 'hours.per.week']]
train_scaler_col = scaler.fit_transform(train_col_sacle)
train_scaler_col = pd.DataFrame(train_scaler_col, columns=train_col_sacle.columns)
df['age'] = train_scaler_col['age']
df['fnlwgt'] = train_scaler_col['fnlwgt']
df['education.num'] = train_scaler_col['education.num']
df['hours.per.week'] = train_scaler_col['hours.per.week']
```

In [111]:

#Data is now divided in independent and dependent.

Creating a train test split

In [112]:

```
from sklearn.model_selection import train_test_split
```

In [113]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.30, random_state=11)
```

In [114]:

```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("Y_train shape:", Y_train.shape)
print("Y_test shape:", Y_test.shape)
```

```
X_train shape: (22708, 14)
X_test shape: (9732, 14)
Y_train shape: (22708,)
Y_test shape: (9732,)
```

In [115]:

#Our data set divided into train and test. Now we will continue with model building.

Data Modelling

Logistic Regression

In [116]:

```
from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression(random_state=42)
```

In [117]:

```
log_reg.fit(X_train, Y_train)
```

Out[117]:

```
LogisticRegression(random_state=42)
```

In [118]:

```
Y_pred_log_reg = log_reg.predict(X_test)
```

KNN Classifier

In [119]:

```
from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier()
```

In [120]:

```
knn.fit(X_train, Y_train)
```

Out[120]:

```
KNeighborsClassifier()
```

In [121]:

```
Y_pred_knn = knn.predict(X_test)
```

Decision Tress

In [122]:

```
from sklearn.tree import DecisionTreeClassifier  
dec_tree = DecisionTreeClassifier(random_state=42)
```

In [123]:

```
dec_tree.fit(X_train, Y_train)
```

Out[123]:

```
DecisionTreeClassifier(random_state=42)
```

In [124]:

```
Y_pred_dec_tree = dec_tree.predict(X_test)
```

Random Forest Classifier

In [125]:

```
from sklearn.ensemble import RandomForestClassifier  
ran_for = RandomForestClassifier(random_state=123)
```

In [126]:

```
ran_for.fit(X_train, Y_train)
```

Out[126]:

```
RandomForestClassifier(random_state=123)
```

In [127]:

```
Y_pred_ran_for = ran_for.predict(X_test)
```

Support Vector Classifier

In [128]:

```
from sklearn.svm import SVC
svc = SVC(random_state=42)
```

In [129]:

```
svc.fit(X_train, Y_train)
```

Out[129]:

```
SVC(random_state=42)
```

In [130]:

```
Y_pred_svc = svc.predict(X_test)
```

Model Evaluation

In [131]:

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
```

In [132]:

```
print('Logistic Regression:')
print('Accuracy score:', round(accuracy_score(Y_test, Y_pred_log_reg) * 100, 2))
print('F1 score:', round(f1_score(Y_test, Y_pred_log_reg) * 100, 2))
```

Logistic Regression:
Accuracy score: 79.47
F1 score: 39.2

In [133]:

```
print('KNN Classifier:')
print('Accuracy score:', round(accuracy_score(Y_test, Y_pred_knn) * 100, 2))
print('F1 score:', round(f1_score(Y_test, Y_pred_knn) * 100, 2))
```

KNN Classifier:
Accuracy score: 77.36
F1 score: 41.55

In [134]:

```
print('Decision Tree Classifier:')
print('Accuracy score:', round(accuracy_score(Y_test, Y_pred_dec_tree) * 100, 2))
print('F1 score:', round(f1_score(Y_test, Y_pred_dec_tree) * 100, 2))
```

Decision Tree Classifier:
Accuracy score: 80.59
F1 score: 61.24

In [135]:

```
print('Random Forest Classifier:')
print('Accuracy score:', round(accuracy_score(Y_test, Y_pred_ran_for) * 100, 2))
print('F1 score:', round(f1_score(Y_test, Y_pred_ran_for) * 100, 2))
```

Random Forest Classifier:
Accuracy score: 85.9
F1 score: 68.63

In [136]:

```
print('Support Vector Classifier:')
print('Accuracy score:', round(accuracy_score(Y_test, Y_pred_svc) * 100, 2))
print('F1 score:', round(f1_score(Y_test, Y_pred_svc) * 100, 2))
```

Support Vector Classifier:
Accuracy score: 78.88
F1 score: 26.48

In [137]:

```
#From the above Model building outcomes it can be analyse that Random Forest Classifier & Decision Tree Classifier  
#the best models with best F1 score and Accuracy score.
```

Hyperparameter Tuning

In [138]:

```
from sklearn.model_selection import RandomizedSearchCV
```

In [139]:

```
n_estimators = [int(x) for x in np.linspace(start=40, stop=150, num=15)]
max_depth = [int(x) for x in np.linspace(40, 150, num=15)]
```

In [140]:

```
param_dist = {
    'n_estimators': n_estimators,
    'max_depth': max_depth,
}
```

In [141]:

```
rf_tuned = RandomForestClassifier(random_state=42)
```

In [142]:

```
rf_cv = RandomizedSearchCV(
    estimator=rf_tuned, param_distributions=param_dist, cv=5, random_state=42)
```

In [143]:

```
rf_cv.fit(X_train, Y_train)
```

Out[143]:

```
RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(random_state=42),
    param_distributions={'max_depth': [40, 47, 55, 63, 71, 79,
    87, 95, 102, 110, 118,
    126, 134, 142, 150],
    'n_estimators': [40, 47, 55, 63, 71, 79,
    87, 95, 102, 110, 118,
    126, 134, 142, 150]},
    random_state=42)
```

In [144]:

```
rf_cv.best_score_
```

Out[144]:

```
0.8561741735434409
```

In [145]:

```
rf_cv.best_params_
```

Out[145]:

```
{'n_estimators': 126, 'max_depth': 79}
```

In [146]:

```
rf_best = RandomForestClassifier(  
    max_depth=102, n_estimators=142, random_state=123)
```

In [147]:

```
rf_best.fit(X_train, Y_train)
```

Out[147]:

```
RandomForestClassifier(max_depth=102, n_estimators=142, random_state=123)
```

In [148]:

```
Y_pred_rf_best = rf_best.predict(X_test)
```


In [149]:

```
print('Random Forest Classifier:')  
print('Accuracy score:', round(accuracy_score(Y_test, Y_pred_rf_best) * 100, 2))  
print('F1 score:', round(f1_score(Y_test, Y_pred_rf_best) * 100, 2))
```

```
Random Forest Classifier:  
Accuracy score: 85.98  
F1 score: 68.87
```

In [150]:

```
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(Y_pred_rf_best, Y_test)  
plt.style.use('default')  
sns.heatmap(cm, annot=True, fmt='d', cmap='YlGnBu')  
plt.xlabel('Predicted classes')  
plt.ylabel('Actual classes')  
plt.savefig('heatmap.png')  
plt.show()
```



In [151]:

```
#Interpretation
```

```
#Y-axis represents the actual classes
```

```
#X-axis represents the predicted classes
```

```
#6859 times when the model correctly predicted 0 when the actual class was 0
```

```
#451 times the model predicted 0 when the actual class was 1
```

```
#913 times the model predicted 1 when the actual class was 0
```

```
#1509 times the model correctly predicted 1 when the actual class was 1
```

In [152]:

```
from sklearn.metrics import classification_report  
print(classification_report(Y_test, Y_pred_rf_best))
```

	precision	recall	f1-score	support
0	0.88	0.94	0.91	7310
1	0.77	0.62	0.69	2422
accuracy			0.86	9732
macro avg	0.83	0.78	0.80	9732
weighted avg	0.85	0.86	0.85	9732

In [153]:

```
#In this project, we build various models like
```

```
# logistic regression
```

```
# knn classifier
```

```
# support vector classifier
```

```
# decision tree classifier
```

```
# random forest classifier
```

```
#A hyperparameter tuned random forest classifier gives the highest accuracy score of 85.98 and f1 score of 68
```

Other method of hyper parameter tuning

In [*]:

```
from sklearn.model_selection import GridSearchCV
n_estimators_List = [40, 47, 55, 63, 71, 79, 87, 95, 102, 110, 118, 126, 134, 142, 150]
max_features_List = [40, 47, 55, 63, 71, 79, 87, 95, 102, 110, 118, 126, 134, 142, 150]
min_samples_leaf_List = [5, 10, 25, 50, 30, 35, 40, 75, 85, 105, 110, 125, 130, 145, 150]

my_param_grid = {'n_estimators': n_estimators_List,
                  'max_features': max_features_List,
                  'min_samples_leaf': min_samples_leaf_List}

Grid_Search_Model = GridSearchCV(estimator = RandomForestClassifier(random_state=123),
                                 param_grid=my_param_grid, scoring='accuracy', cv=3).fit(X_train, Y_train) # param_grid is a dictionary

Model_Validation_Df4 = pd.DataFrame.from_dict(Grid_Search_Model.cv_results_)
# Grid_Search_Model.cv_results_

# Based on the selected hyperparameters, you should build a final model on the COMPLETE training data (trainX, Y_train)
RF_Final = RandomForestClassifier(random_state = 123, n_estimators = 75,
                                 max_features = 9, min_samples_leaf = 5).fit(X_train, Y_train)

Test_Pred = RF_Final.predict(X_test)

# Confusion Matrix
Confusion_Mat = pd.crosstab(Y_test, Test_Pred) # R, C format (Actual = testY, Predicted = Test_Pred)
Confusion_Mat

# Validation on Testset
print(classification_report(Y_test, Test_Pred)) # Actual, Predicted
```

In []: